

Review of AI Ethics Interview with Reid Blackman

This interview features Reid Blackman, a philosopher, ethicist, and founder of the AI ethical risk consultancy Virtue, discussing practical approaches to responsible artificial intelligence. The conversation offered actionable insights for students and practitioners entering technology fields. Blackman draws on his academic background in philosophy, entrepreneurial experience, and consulting work with Fortune 500 organizations to present a framework for addressing ethical risks in AI development and deployment.

Core Ethical Risks Inherent to AI

Blackman argues that AI ethics differs from general applied ethics because the architecture of machine learning systems systematically increases the probability of specific harms. He identifies three primary risks. First, bias and discrimination can emerge not from malicious intent but from technical choices regarding training data, variable weighting, decision thresholds, and objective functions. Second, the black box problem means complex models produce outputs through mathematical patterns too intricate for human interpretation, creating accountability challenges in high stakes decisions. Third, privacy violations arise through extensive data collection, inference of sensitive information from seemingly neutral data, and surveillance applications like facial recognition.

Blackman contrasts AI with simple tools such as screwdrivers. While any technology can be misused, AI's operational nature elevates the likelihood of these particular ethical failures, requiring proactive mitigation rather than reactive responses.

AI for Not Bad

A central contribution is Blackman's recommendation that organizations prioritize avoiding harm over pursuing aspirational positive social impact. He observed that AI for Good rhetoric often fails to influence business leaders focused on operational realities. Instead, he proposes an AI for Not Bad approach: pursuing legitimate business objectives while ensuring the means do not unacceptably risk wronging society. Reframing ethics as risk management aligned with core business interests including brand protection, legal compliance, and stakeholder trust. The framework acknowledges that companies need not achieve moral excellence to operate responsibly, but they must avoid causing discriminatory, unexplainable, or privacy invasive outcomes.

Practical Guidance on Bias Mitigation

Blackman emphasizes that bias has multiple sources beyond training data quality. Effective mitigation requires proactive consideration at the concept phase, asking whether a proposed use case or data source could produce discriminatory impacts before model development begins. Assessment must continue throughout the AI lifecycle because changes at any stage, such as removing a dataset for privacy reasons, can reintroduce bias risks. Solutions must be context sensitive; representative data may help in one scenario while adjusting

thresholds or objective functions may be needed in another. Organizations should document their bias assessment processes and assign clear responsibilities using frameworks like RACI matrices to ensure accountability.

Distinguishing Explainability from Transparency

Blackman explains that these explainability and transparency serve different purposes and should not be conflated. Transparency concerns communication with stakeholders about what AI is being used for, for what purpose, and with what data, while explainability concerns technical understanding of how a model arrived at a specific output. Organizations can be transparent about using a black box model or opaque about using an interpretable one. The appropriate balance depends on context and stakes. High stakes applications such as healthcare or lending demand greater explainability than low stakes applications like fraud detection flagging for human review. Blackman tells us that there are two approaches: inherently interpretable models for contexts requiring clear reasoning, and post hoc explanation techniques for complex models, with the caveat that simplifications must remain accurate enough to support trustworthy decisions.

Building Effective AI Governance

Blackman critiques vague AI ethics statements as ineffective unless operationalized. Effective governance requires action oriented commitments that translate values into specific guardrails, such as "We will never sell individual consumer data to third parties." Oversight must be cross functional because ethical decisions involve qualitative judgments beyond technical expertise, requiring input from legal professionals, civil rights experts, ethicists, and business leaders. Accountability requires measurable tracking of whether promised actions like bias assessments and staff training are actually implemented. Governance must extend enterprise wide beyond in house data science teams to procurement and business units, as externally sourced AI tools can introduce unvetted risks.

AI and Blockchain Considerations

Blackman acknowledges the appeal of democratizing AI through no code or low code tools that enable business professionals to prototype solutions. However, he warns that technical and ethical expertise cannot be fully automated. Novice practitioners may inadvertently create harmful models without proper training or oversight. Organizations pursuing this approach must invest in robust training, mentorship, and expert review processes before deployment.

Regarding blockchain, Blackman challenges the narrative that decentralization eliminates governance needs. While processing and storage may be distributed, control often remains concentrated among developers. Assuming that code alone ensures trust can obscure accountability and enable harmful uses.

Relevance for Graduating Students

This interview offers valuable insights for students entering technology fields. Blackman's emphasis on cross functional collaboration shows the importance of complementing

technical skills with ethical reasoning, communication abilities, and business context awareness. His practical frameworks, such as lifecycle bias assessment and action oriented ethics statements, provide concrete tools applicable in internships and careers. In addition, Blackman's career trajectory demonstrates how deep disciplinary expertise can translate into high impact industry roles when combined with business acumen. His transition from philosophy professor to ethics consultant illustrates that ethical reasoning is critical to sustainable innovation.

Conclusion

Reid Blackman's interview provides a comprehensive yet accessible overview of key ethical challenges in artificial intelligence. By grounding abstract principles in practical governance strategies, he offers a roadmap for people seeking to deploy AI responsibly. His emphasis on proactive risk mitigation, context sensitive decision making, and measurable accountability provides actionable guidance for people working in the technology industry. For students preparing to enter technology fields, the conversation reinforces that ethical considerations are not obstacles to innovation but essential components of building trustworthy, sustainable systems. The interview makes complex ethical concepts accessible, allowing university students and members of the community to learn about how to build and use AI more ethically.